PHASE P

Hard revert to J.1.6 + Market Mode Detector build

POLYMARKET 5-minute BTC up/down bot
Reverted: every Phase L/M/O bot-runtime feature (code + env)
Restored: J.1.6 peak config (Kelly 0.15× + bankroll replay + cap=1)
Branch: **claude/phase-P-hard-reset** @ **53e08c9** (cut from 83fe119)
Session 1 close: **2026-05-17 21:37:34 UTC**

Bot LIVE on J.1.6 baseline at $1,000 bankroll — MainPID 1019607

Source-of-truth prompt: ~/Desktop/phase-P-hard-reset-2026-05-18.txt
Authored by operator after Phase O Session 1 disappointment.

Contents

§1	What we're reverting and why — operator's framing
§2	Phase N data summary — the 5-day forensic evidence
§3	J.1.6 baseline restored — exact env + branch commit
§4	Market Mode Detector (MMD) — Sessions 2-3 thesis
§5	Bot state confirmation — startup log + dashboard at $1k anchor
§6	Three-session phased roadmap + acceptance criteria
§7	Forensic data preservation — where the 5-day data lives

This kickoff PDF is the Session-1 deliverable for Phase P. It does NOT argue against the revert. It documents what was reverted, the data that justifies the verdicts, the exact baseline now LIVE, and the path to Phase Q.

§1 What we're reverting and why

Operator's verbatim framing (source-of-truth prompt at `~/Desktop/phase-P-hard-reset-2026-05-18.txt`):

> *"I'm really disappointed at the results that we currently have. We're going to revert back to pre peak. The pre peak — the last ship was J.1.6, Kelly 0.15× plus replay. So yeah, just let's revert back to pre peak."*
>
> *"Use GitHub. Make sure the VPS, the Mac directory, and everything is aligned together with the same version. We need you to revert. So just make a plan for that because we have shipped so many fucking features based on emotions post peak. ... The bot is in a really nice stage. It can't even make a single trade now. Fully revert to pre peak now."*
>
> *"I think we're missing one thing only, and I think it's the regime classifier. Let's not call it regime classifier. ... find this one GitHub repo that when we turn on and integrate into our bot, it can detect when Bitcoin is going up, when Bitcoin is going down, and when Bitcoin is choppy or mixed."*
>
> *"Make sure to restart the bot, turn it on with one thousand dollar bankroll. Make sure to reset the chart, the cumulative PnL chart, and make sure to reset the decision tape. Just reset the full dashboard. Make a full plan. This is the last chance we have."*

The story arc this revert is responding to:

Phase	Date	What
Phase G	2026-05-12	Zero-edge audit. 50% WR / –$19 net on n=633.
Phase H	2026-05-13	v2.1 priors + Kelly + cap=1 + H.2.12 path-A block_up.
Phase J	2026-05-14	J.1.5 CH3 fold + J.2 hour blocklist + J.1.6 Kelly 0.15× + replay. ← THE PEAK CONFIG.
Phase K	2026-05-16	4-day soak. Peak +$1,235 @ 05-15 15:40Z. Day-4 drawdown –$1,294 in 18h.
Phase L	2026-05-16	SUPERSEDED same-day by Phase M.
Phase M	2026-05-16/17	5 layers in 22h: L.5 + M.1a v0 + M.1a v1 + M.5 + M.2 v0.
Phase N	2026-05-17	Forensic PDF (22 pp). M.1a v1 FALSIFIED on n=755. M.5 design CORRECT but shipped 3 days late. SURV
Phase O	2026-05-17	Session 1 remediation. 3 emergency mid-session pivots ($1k→$3k, L.5 ×3, M.5 lookback). 1 fire / 5h = LOS
Phase P	2026-05-18	THIS PHASE — strip post-peak ware, restore J.1.6, add ONE new thing (MMD) on data-grounded process.

J.1.6 (the revert target) was live at the peak *and* through the post-peak drawdown. The operator knows this. The revert isn't a magic fix — it's a clean baseline from which to add ONE measured ingredient with backtest + shadow-log discipline.

§2 Phase N data summary — what the 5 days proved

Phase N forensic audit (pdfs/PHASE_N_FORENSIC_AUDIT.pdf, 22 pp, 14 charts) audited n=755 settled fires across 4.2 days (100% strategy=spot_aware_v2_1, mode=kelly_multi). It is the data backbone for every KEEP/DROP verdict below.

§2.1 Headline findings (Phase N §3-§5)

Finding	What the data shows
SURV cell (tier 30-50% × |Δbps|≥5)	Pre-peak WR 60.7%, PnL +$1,431 → Post-peak WR 46.5%, PnL –$1,161. <b>Dominant bleeder.</
Tier 70-85% Path A (H.2.12 block_up)	n=237, 0 UP fires (gate works); DOWN-only WR 74.3% vs ~77.5% break-even. –$368 5-day net. Sl
CH3 cell (J.1.5 fold)	0 fires post-fold across n=755. Fold works as designed.
M.1a v1 chop-abstain	FALSIFIED. Annotated n=83: chop n=23 WR 60.9% +$31 vs mixed n=58 WR 44.8% –$296. Falsifie
M.5 path-A regime wrapper	Design CORRECT but shipped 3 days late. Idealized simulator: +$352.92 counterfactual avoided. R
M.2 v0 orderbook imbalance shadow	Lived 45 min; n=3, all too_thin, all lost. Un-evaluable.
Restart erosion	19 likely bot restarts in 5 days (one 177-min outage). Per-restart cold-start cost ~$22 to M.5's counte

§2.2 Why this justifies the revert (not just M.5 ship-late fix)

Phase O Session 1 already implemented Phase N's #1 recommendation (disable M.1a v1) but layered on additional remediations (SURV canary + bankroll override $1k→$3k + L.5 ×3 + M.5 lookback tune). The 1 post-restart fire = LOSS –$163 didn't refute the design — but the operator concluded that the Phase M/O *stack* had become un-tractable and that compounding fixes on top wasn't going to clear the picture.

Hard revert achieves: (1) clean baseline for measuring any new ingredient in isolation; (2) restoring decision #10 ($1k bankroll), keeping bet sizes in the $25-150 band where Phase J variance was characterized; (3) clear decoupling between "the old defenses fight each other" vs "the underlying strategy has degraded".

> Trade-off documented: J.1.6 is the same config that produced the $1,235 peak *and* the $1,294 drawdown. The revert does not promise a different outcome at $1k stakes. It promises a clean slate against which the MMD (§4) can be measured.

§3 J.1.6 baseline restored — exact env + branch

§3.1 Branch state

Branch	claude/phase-P-hard-reset
HEAD	53e08c9 (chore(phase-P): preserve Phase N forensic data)
Cut from	83fe119 — docs(phase-J): design lock for J.3 + J.7 (docs-only, last commit before L.5 shipped)
Author of base	Hamad, 2026-05-13 11:28:23 +0400 (07:28 UTC)
Commits dropped vs jovial-hermann@e3133e2	16 (all Phase L docs + M.* + O.*)
Files preserved on new branch	Phase N PDF + builder + scripts + 14 PNG charts + runbooks + ledgers + 4 session notes

§3.2 KEEP-LIST (the J.1.6 strategy surface)

Env var	Value	Origin
STRATEGY	spot_aware_v2_1	H.2 v2.1
KELLY_MULT	0.15	J.1.6 B — Tier S1
MULTI_FIRE_CAP	1	H.5 throttle
BANKROLL_REPLAY_ENABLED	true	J.1.6 A — Kelly compounds off bankroll
PMB_STARTING_BANKROLL	1000	decision #10 restored from $3k Phase O override
SPOT_AWARE_TIER_RELAX_P_WIN_FALLBACK	0.0	J.1.5 fold — CH3 cell dead
SPOT_AWARE_TIER_RELAX_PRICE_FLOOR	0.42	J.1 — H.3 audit n=111 boundary
SPOT_AWARE_HOUR_BLOCKLIST_UTC	22,23,0	J.2 — catastrophic hours × 70-85%
SPOT_AWARE_HOUR_BLOCKLIST_TIERS	70-85%	J.2
SPOT_AWARE_BLOCK_UP_IN_TIERS	70-85%	H.2.12 path-A — block UP in tier 70-85%
INSTANT_SETTLE_ENABLED	1	cross-phase (latency)
INSTANT_SETTLE_GAMMA_TIMEOUT_S	30	H.2.10
POST_HOC_GAMMA_RETRY_ENABLED	1	cross-phase (settle truth)

§3.3 DROP-LIST (git-reverted + .envrc stripped)

Feature	How it was removed
L.5 daily-PnL auto-halt	design + src/_live/auto_halt.py + DAILY_PNL_* env + auto-resume cron
M.1a v0 + v1 regime classifier	design + spot_aware_v2_1.py gate + REGIME_M1A_* + REGIME_CLASSIFIER_V2_SHADOW_
M.5 path-A regime wrapper	design + path_a_wr_metrics.py + REGIME_M5_* env
M.2 v0 orderbook imbalance shadow	design + spot_aware_v2_1.py annotation + ORDERBOOK_IMBALANCE_* env
O.1 SURV-cell shift canary	design + surv_canary_metrics.py + SURV_CANARY_* env
Dashboard PMB env passthrough	src/.../dashboard/bankroll.py + dashboard_server/state.py mod + systemd drop-in
Phase O bankroll override $1k→$3k	env-only (PMB_STARTING_BANKROLL=3000, BANKROLL_REPLAY_ENABLED=false)
Phase 14 tier-gate env	kept code (no-op default false); stripped TIER_GATE_ENABLED / TIER_GATE_BLOCKLIST from

Feature	How it was removed
Phase 15 regime-gate env	kept code (no-op default false); stripped REGIME_GATE_ENABLED / REGIME_GATE_ABSTAIN

Full inventory with Phase N verdict per feature at `runbooks/research/phase-P-feature-inventory.md`.

§4 Market Mode Detector (MMD) — Sessions 2-3 thesis

§4.1 Why MMD (not "regime classifier")

Operator-directed rename: the previous "regime classifier" name is deprecated. The new name is **Market Mode Detector (MMD)**. The intent is the same shape but a different lineage from the dead M.1a vol-ratio classifier:

Dimension	M.1a (dead) vs MMD (this phase)
Sourced from	M.1a v0/v1: hand-rolled vol-ratio | MMD: battle-tested GitHub TA library
Inputs	M.1a: single vol-ratio | MMD: multi-indicator (ADX + ATR + BB-width + momentum)
Output labels	M.1a: trending / chop / mixed (3) | MMD: UP / DOWN / CHOPPY / MIXED (4) + optional confidence
Validation discipline	M.1a v1: activated on n=23 annotated subset → falsified on n=755 | MMD: validated against Phase N n=755 BI
Ship discipline	M.1a v0 → v1 in 2 days (DARK → ACTIVE) | MMD: Session 2 backtest → Session 3 SHADOW only → Day-5 a

§4.2 What MMD outputs (target spec)

Given the spot tick stream (Binance + Coinbase WebSocket already running) plus a rolling window, output a label:

Label	Meaning
UP	sustained uptrend
DOWN	sustained downtrend
CHOPPY	high-frequency low-amplitude oscillation
MIXED	neither clearly trending nor clearly choppy

Plus an optional confidence $\in$ [0,1] and a forward "what's likely next" projection (operator framing: "kind of know or predict what's gonna happen next").

§4.3 Candidate repos Session 2 will evaluate

Repo	Why on the shortlist
bukosabino/ta	pure-Python TA (ADX/ATR/BB/momentum). Battle-tested, MIT. Likely the primary candidate.
twopirllc/pandas-ta	large TA library on pandas. MIT.
hudson-and-thames/mlfinlab	ML-for-finance regime-switching modules. MIT.
hmmlearn	Hidden Markov regime models. BSD.
GitHub search	top 5 repos ≥100 stars + ≥90d recent activity matching "bitcoin regime detection" / "cry

Evaluation criteria: tick-input or trivial bar-aggregation; structured output (label + confidence); license MIT/Apache/BSD (decision #8); backtestable on Phase N n=755 forensic data; dep footprint < 50 MB + no uv.lock conflict.

Session 2 will extend scripts/research/phase_n_forensic_audit.py with a `section_p_mmd_backtest()` that classifies each of the 755 fires, joins to outcome, computes per-label WR + $/tr table, and scores reversal-prediction TPR via a ROC sweep. A USEFUL MMD shows ≥ 10pp WR delta between best and worst label on n=755 AND TPR ≥ 0.6 on the reversal-prediction task.

§5 Bot state confirmation — $1k anchor proven

§5.1 systemd state (VPS, 2026-05-17 21:37:34 UTC)

bot-staging.service	active / running
MainPID	1019607
ExecMainStartTimestamp	Sun 2026-05-17 21:37:28 UTC
dashboard-server.service	active (restart with bankroll.conf drop-in removed)
state/halt.sentinel	absent
VPS HEAD	53e08c903f4d2a3e53b8da63aed35e2184c8e236
VPS branch	claude/phase-P-hard-reset

§5.2 Bot startup log (canonical J.1.6 signature)

```
active strategy: spot_aware_v2_1
spot_aware_v2_1: kelly_mult=0.15 max_bet_abs=$500 multi_fire_cap=1
tier_relax_p_win_fallback=None tier_relax_price_floor=0.420
block_up_in_tiers=('70-85%',) block_down_in_tiers=()
hour_blocklist_utc=(22, 23, 0) hour_blocklist_tiers=('70-85%',)
instant-settle ENABLED: gamma poll timeout=30.0s interval=5.0s
post-hoc gamma retry ENABLED: initial_delay=300.0s interval=30.0s
spot_aware_v2_1: bankroll_getter wired (start=$1000.00, replay=True, cap=1)
bankroll replay: starting=$1000.00 current=$1000.00
orphan recovery: 0 orphans found
```

Absent (intentional): L.5 / M.5 / M.1a / SURV-canary / orderbook-imbalance / MMD lines.

§5.3 Dashboard /api/state at session close

mode.label	WAITING
bankroll.starting_usd	$1000.0
bankroll.current_usd	$1000.0
bankroll.delta_usd	$0.0
today.n_settled	0
decisions[]	0 entries
cumulative_pnl_series	0 pts
scenarios[].n_settled (sample)	0

Dashboard fully reset per operator directive ("reset the chart, the cumulative PnL chart, and reset the decision tape. Just reset the full dashboard"). Systemd drop-in `/etc/systemd/system/dashboard-server.service.d/bankroll.conf` (Phase O env passthrough) was deleted; daemon-reloaded; dashboard restarted.

§6 Three-session phased roadmap

Session	When	Status	Outcome
Session 1	2026-05-17 21:37 UTC	DONE	Hard revert + dashboard reset + this PDF. Bot LIVE on J.1.6 baseline. Vault + mer
Session 2	T+1-3 days (operator GO)	PENDING	MMD candidate research (top-2 from shortlist). Extend phase_n_forensic_audit.py
Session 3	T+5-7 days after Session 2	PENDING	Implement MMD per locked design. ≥12 new tests + mypy strict + ruff + import-linte
Phase Q decision	Day 5 from Session 3 ship	—	MMD ≥10pp WR delta → promote to ACTIVE gate (consider also re-introducing M

§6.1 Session 1 acceptance — all green

- ✓ Mac main repo HEAD == VPS HEAD == origin/claude/phase-P-hard-reset
- ✓ bot-staging.service active on new HEAD; J.1.6 baseline in startup log
- ✓ Dashboard /api/state: $1,000 starting / $1,000 current / $0 today PnL / empty decisions[] / empty cumulative_pnl_series
- ✓ state/decisions/2026-05-18.jsonl exists (empty)
- ✓ state/decisions{,-resolved}/archived/phase-P-reset/*.jsonl preserves 5-day data
- ✓ Dashboard systemd drop-in bankroll.conf REMOVED
- ✓ state/halt.sentinel absent
- ✓ pdfs/PHASE_P_KICKOFF.pdf rendered (this file)
- ✓ Vault + memory + MEMORY.md updated; Phase O marked SUPERSEDED
- ✓ No Phase M/O bot-runtime code in src/ on the new branch
- ✓ Forensic data (Phase N PDF + audit script + charts) survives on new branch

§6.2 Hard DO-NOTs (operator-stated)

- × Do not re-add any Phase M/O bot-runtime feature during Session 2 or 3.
- × Do not change the bankroll from $1,000 (decision #10 restored).
- × Do not activate MMD on small-n during Session 3 — shadow first, always.
- × Do not re-introduce M.5 / SURV canary in Phase P — wait for the Phase Q decision after Day-5 MMD re-soak data.
- × Do not delete the forensic data — archive only.
- × Do not include preachy honest-pushback sections in the PDFs that argue against the revert.

§7 Forensic data preservation map

§7.1 On the new branch (cherry-picked from e3133e6)

Path	What
pdfs/PHASE_N_FORENSIC_AUDIT.pdf	22 pp, 14 charts, 1.3 MB
pdfs/build_phase_n_forensic_audit_pdf.py	reportlab builder
scripts/research/phase_n_forensic_audit.py	n=755 audit + simulators (380 LOC + Session 2 extensions)
scripts/research/phase_n_charts.py	matplotlib chart pipeline (~400 LOC)
runbooks/research/phase-N-forensic-output-2026-05-17.txt	raw audit output
runbooks/research/phase-N-forensic-summary.md	operator briefing
state/research/phase-N-charts/*.png	14 PNGs (cumulative PnL, bankroll trajectory, daily PnL bar, per-tier PnL/WR, SU
state/phases/phase-N-forensic-audit.md	Phase N phase ledger
state/phases/phase-O-remediation.md	Phase O phase ledger (historical)
state/sessions/2026-05-17-phase-N-kickoff.md	Phase N Session 1
state/sessions/2026-05-17-phase-N-session-2-pdf-study.md	Phase N Session 2 (PDF study)
state/sessions/2026-05-17-phase-O-session-1-bot-resumed.md	Phase O Session 1
state/sessions/2026-05-18-phase-O-session-2-prep.md	Phase O Session 2 prep (superseded by Phase P)

§7.2 On the VPS (archived runtime data)

Path	Size / note
state/decisions/archived/phase-P-reset/2026-05-13.jsonl	240,427 bytes
state/decisions/archived/phase-P-reset/2026-05-14.jsonl	236,536 bytes
state/decisions/archived/phase-P-reset/2026-05-15.jsonl	205,346 bytes (peak day)
state/decisions/archived/phase-P-reset/2026-05-16.jsonl	233,606 bytes (drawdown day)
state/decisions/archived/phase-P-reset/2026-05-17.jsonl	70,993 bytes (partial — Phase O day)
state/decisions-resolved/archived/phase-P-reset/2026-05-{13..17}.jsonl	240K / 208K / 237K / 72K (post-hoc gamma-truth resolutions)
/tmp/phaseN_data/state/decisions{,-resolved}/	Original Phase N tarball (re-pull command in PHASE_N audit script)
.envrc.bak.phase-P-revert-2026-05-18	Pre-revert .envrc (8,404 bytes, contains all Phase M/O env vars for reference
.envrc.bak.* (older backups)	Phase F / Phase H / Phase O / etc. — kept for historical reference

§7.3 What Session 2 MMD backtest will use

The Session-2 MMD backtest extension to `scripts/research/phase_n_forensic_audit.py` will: (1) load the JSONL files from `state/decisions{,-resolved}/archived/phase-P-reset/`, (2) for each of the 755 fires reconstruct ~30 minutes of spot trajectory from snapshot.spot_now + oracle_last, (3) classify each moment with the candidate MMD, (4) join to per-fire outcome, (5) compute per-label WR + $/tr + n table + reversal-prediction ROC.

§8 Operator's bottom line + Phase Q decision matrix

§8.1 Operator's bottom line (verbatim)

> *"The bot is in a really nice stage. It can't even make a single trade now. Fully revert to pre peak now. Make a fully phased plan. This is the last chance we have."*

— operator, 2026-05-18

"The last chance" framing is the operating frame for Phase P, Q and beyond. The clock isn't capital (J.1.6 + revert leaves the bankroll intact at $1k); it's the operator's trust in the design loop. The remedy is process discipline: shadow before active, n=755 backtest before ship, ONE new thing per phase. Phase P is the demonstration that we can hold to that.

§8.2 Phase Q decision matrix (set NOW so we don't re-decide later)

Phase Q signal	Action
≥10pp WR delta between best and worst MMD label on n=755 backtest AND shadow-validated	PROMOTE MMD to ACTIVE gate (MMD_MODE=gate, MMD_LABELS_T
5-10pp WR delta	KEEP MMD in SHADOW; iterate on indicator parameters (lookback / threshold). R
≤5pp WR delta	SHELVE MMD; reconsider the strategy class entirely. Path to Phase R: either pivo
MMD fails to compute (classifier exception OR ambiguous data)	Pass-through fail-safe (already in design): do not gate. Telemetry counter tracked.

§8.3 What this whole phase unlocks

Milestone	Outcome
After Session 1 (NOW)	Bot LIVE on clean J.1.6 at $1k. Mac + VPS + GitHub aligned. Dashboard fully reset. 5-day forens
After Session 2	MMD candidate picked + backtested on n=755. Integration design locked. PHASE_P_SESSION2
After Session 3	MMD shipped + shadow-annotating fires (mmd=<label> in rationale). Bot still on J.1.6 (no firi
After Phase P CLOSED (~Day 5 from Session 3 ship)	Empirical verdict on MMD as edge source. ~9 days of total data (Phase N's 5 + 5 new). Phase Q
Phase Q decision	KEEP J.1.6 + MMD as Phase R baseline, OR pivot strategy class, OR reconsider entire approach

§9 Glossary + version manifest

§9.1 Acronym + named-element glossary

Term	Meaning
J.1.6	Phase J Session 1.6 (2026-05-14 18:36 UTC) — Kelly 0.15× Tier S1 + BANKROLL_REPLAY_ENABLED=tru
J.1.5 fold	Phase J Session 1.5 (2026-05-14 07:22 UTC) — SPOT_AWARE_TIER_RELAX_P_WIN_FALLBACK 0.48 →
H.2.12 path A	Phase H 2.12 (2026-05-13 00:56 UTC, commit c746db3) — SPOT_AWARE_BLOCK_UP_IN_TIERS=70-85%
CH3 cell	Conservative-Hedge cell: tier 30-50% with |Δbps|<5. Phase G winning cell on n=83; turned out to lose live.
SURV cell	Tier 30-50% × |Δbps|≥5. Dominant bleeder identified by Phase N (60.7% → 46.5% WR collapse post-peak).
L.5	Phase M Session 2 (2026-05-16) — daily-PnL auto-halt at -$30 gate, check 23:50 UTC. STRIPPED in Phase
M.1a v0 / v1	Phase M Session 4/5 — FLAT-vs-TREND regime classifier (vol-ratio). v0 = SHADOW, v1 = ACTIVE chop-abs
M.5	Phase M Session 6 — path-A regime-aware wrapper with rolling-WR threshold. Design correct; shipped 3 day
M.2 v0	Phase M Session 7 — orderbook imbalance shadow-log. n=3 un-evaluable; STRIPPED in Phase P.
O.1 SURV canary	Phase O Session 1 — SURV-cell shift detector with rolling-WR margin. n=1 post-restart fire; STRIPPED in Ph
MMD	Market Mode Detector. Phase P Session 2-3 deliverable. 4 labels (UP/DOWN/CHOPPY/MIXED) + optional co
Decision #10	Locked decision: $1,000 starting bankroll. Overridden in Phase O ($1k → $3k); RESTORED in Phase P.

§9.2 Version manifest (Session 1 close)

Branch	claude/phase-P-hard-reset
HEAD	53e08c9 (chore(phase-P): preserve Phase N forensic data)
Base	83fe119 (docs(phase-J): design lock for J.3 + J.7)
GitHub remote	https://github.com/FalconAEZ/polymarket-5m-bot
VPS host:port	ops@108.132.51.116:2222 (Dublin Lightsail, eu-west-1)
VPS service	bot-staging.service MainPID 1019607 (active since 2026-05-17 21:37:28 UTC)
VPS .envrc backup	.envrc.bak.phase-P-revert-2026-05-18 (8,404 bytes)
Bot startup signature	kelly_mult=0.15 / cap=1 / tier_relax_price_floor=0.420 / block_up_in_tiers=('70-85%',) / hour_blockli
Dashboard URL	http://127.0.0.1:8505 (SSH tunnel only; never public)
This PDF builder	pdfs/build_phase_p_kickoff_pdf.py
This PDF output	pdfs/PHASE_P_KICKOFF.pdf
Phase ledger	state/phases/phase-P-hard-reset.md
Session note	state/sessions/2026-05-18-phase-P-session-1-revert.md
Memory checkpoint	~/.claude/projects/-Users-hamad-Desktop-PolyMarket-5m/memory/checkpoint_phase_P.md
Feature inventory	runbooks/research/phase-P-feature-inventory.md
Source-of-truth prompt	~/Desktop/phase-P-hard-reset-2026-05-18.txt

End of Phase P Session 1 kickoff. Bot is LIVE on J.1.6 baseline at $1,000 bankroll. Next: operator GO for Session 2 MMD research.